

1. Importing the dependencies

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
```

2. Data Loading and Understanding

```
#load the csv data to a pandas dataframe
df = pd.read_csv("/content/WA_Fn-UseC_-Telco-Customer-Churn.csv.xls")
```

```
df.shape
```

```
(7043, 21)
```

```
df.head()      #may not show some column... some columns might get truncated
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	

```
pd.set_option("display.max_columns", None)      #will display every column
```

```
df.head(2)      #total charges = monthly charges * tenure(in month)
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null    object  
 1   gender          7043 non-null    object  
 2   SeniorCitizen   7043 non-null    int64  
 3   Partner         7043 non-null    object  
 4   Dependents     7043 non-null    object  
 5   tenure          7043 non-null    int64  
 6   PhoneService    7043 non-null    object  
 7   MultipleLines   7043 non-null    object  
 8   InternetService 7043 non-null   object  
 9   OnlineSecurity  7043 non-null   object  
 10  OnlineBackup    7043 non-null   object  
 11  DeviceProtection 7043 non-null  object  
 12  TechSupport    7043 non-null   object  
 13  StreamingTV    7043 non-null   object  
 14  StreamingMovies 7043 non-null  object  
 15  Contract        7043 non-null   object  
 16  PaperlessBilling 7043 non-null  object  
 17  PaymentMethod   7043 non-null   object  
 18  MonthlyCharges 7043 non-null   float64 
 19  TotalCharges   7043 non-null   object  
 20  Churn          7043 non-null   object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
#dropping customerID column as this is not required for modelling
df = df.drop(columns=["customerID"])
```

```
df.head(2)          #customerID is removed
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	...
0	Female	0	Yes	No	1	No	...
1	Male	0	No	No	34	Yes	...

Next steps:

[Generate code with df](#)

[New interactive sheet](#)

```
df.columns
```

```
Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents',
       'tenure',
       'PhoneService', 'MultipleLines', 'InternetService',
       'OnlineSecurity',
       'OnlineBackup', 'DeviceProtection', 'TechSupport',
       'StreamingTV',
       'StreamingMovies', 'Contract', 'PaperlessBilling',
       'PaymentMethod',
       'MonthlyCharges', 'TotalCharges', 'Churn'],
       dtype='object')
```

```
print(df["gender"].unique())      #categorical
['Female' 'Male']
```

```
print(df["SeniorCitizen"].unique())      #categorical
[0 1]
```

```
#printing the unique values in all the columns / printing only the numerical_features_list = ["tenure", "MonthlyCharges", "TotalCharges"]

for col in df.columns:
    if col not in numerical_features_list:
        print(col, df[col].unique())
        print("-"*50)

gender ['Female' 'Male']
-----
SeniorCitizen [0 1]
-----
Partner ['Yes' 'No']
-----
Dependents ['No' 'Yes']
-----
PhoneService ['No' 'Yes']
-----
MultipleLines ['No phone service' 'No' 'Yes']
-----
InternetService ['DSL' 'Fiber optic' 'No']
-----
OnlineSecurity ['No' 'Yes' 'No internet service']
-----
OnlineBackup ['Yes' 'No' 'No internet service']
-----
DeviceProtection ['No' 'Yes' 'No internet service']
-----
TechSupport ['No' 'Yes' 'No internet service']
-----
StreamingTV ['No' 'Yes' 'No internet service']
-----
StreamingMovies ['No' 'Yes' 'No internet service']
-----
Contract ['Month-to-month' 'One year' 'Two year']
-----
PaperlessBilling ['Yes' 'No']
-----
PaymentMethod ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
  'Credit card (automatic)']
-----
Churn ['No' 'Yes']
```

```
print(df.isnull().sum())      #will count the number of missin  
  
gender          0  
SeniorCitizen  0  
Partner         0  
Dependents     0  
tenure          0  
PhoneService   0  
MultipleLines   0  
InternetService 0  
OnlineSecurity  0  
OnlineBackup    0  
DeviceProtection 0  
TechSupport    0  
StreamingTV    0  
StreamingMovies 0  
Contract        0  
PaperlessBilling 0  
PaymentMethod   0  
MonthlyCharges  0  
TotalCharges    0  
Churn           0  
dtype: int64
```

```
df[df["TotalCharges"]==" "]
```

		gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	StreamingTV	StreamingMovies
488		Female	0	Yes	Yes	0	No	No	No
753		Male	0	No	Yes	0	Yes	Yes	Yes
936		Female	0	Yes	Yes	0	Yes	Yes	Yes
1082		Male	0	Yes	Yes	0	Yes	Yes	Yes
1340		Female	0	Yes	Yes	0	No	No	No
3331		Male	0	Yes	Yes	0	Yes	Yes	Yes
3826		Male	0	Yes	Yes	0	Yes	Yes	Yes
4380		Female	0	Yes	Yes	0	Yes	Yes	Yes
5218		Male	0	Yes	Yes	0	Yes	Yes	Yes
6670		Female	0	Yes	Yes	0	Yes	Yes	Yes
6754		Male	0	No	Yes	0	Yes	Yes	Yes

```
len(df[df["TotalCharges"]==" "])
```

```
11
```

```
df["TotalCharges"] = df["TotalCharges"].replace({" ":" "0.0"})
```

```
df["TotalCharges"] = df["TotalCharges"].astype(float)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   gender            7043 non-null    object  
 1   SeniorCitizen     7043 non-null    int64  
 2   Partner           7043 non-null    object  
 3   Dependents        7043 non-null    object  
 4   tenure            7043 non-null    int64  
 5   PhoneService      7043 non-null    object  
 6   MultipleLines     7043 non-null    object  
 7   InternetService   7043 non-null    object  
 8   OnlineSecurity    7043 non-null    object  
 9   OnlineBackup       7043 non-null    object  
 10  DeviceProtection  7043 non-null    object  
 11  TechSupport       7043 non-null    object  
 12  StreamingTV       7043 non-null    object  
 13  StreamingMovies   7043 non-null    object  
 14  Contract          7043 non-null    object  
 15  PaperlessBilling  7043 non-null    object  
 16  PaymentMethod     7043 non-null    object  
 17  MonthlyCharges   7043 non-null    float64 
 18  TotalCharges      7043 non-null    float64 
 19  Churn             7043 non-null    object  
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB
```

```
# checking the class distribution of target column
print(df["Churn"].value_counts())
```

```
Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

Insights:

1. Customer ID removed as it is not required for modelling
 2. No missing values in the dataset
 3. Missing values in the TotalCharges column were replaced with 0
 4. Class imbalance identified in the target
-
3. Exploratory Data Analysis (EDA)

```
df.shape
```

```
(7043, 20)
```

```
df.columns
```

```
Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents',  
'tenure',  
       'PhoneService', 'MultipleLines', 'InternetService',  
'OnlineSecurity',  
       'OnlineBackup', 'DeviceProtection', 'TechSupport',  
'StreamingTV',  
       'StreamingMovies', 'Contract', 'PaperlessBilling',  
'PaymentMethod',  
       'MonthlyCharges', 'TotalCharges', 'Churn'],  
      dtype='object')
```

```
df.head(2)
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	...
0	Female	0	Yes		No	1	No
1	Male	0	No		No	34	Yes

Next steps:

[Generate code with df](#)[New interactive sheet](#)

```
df.describe()
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692	2279.734304
std	0.368612	24.559481	30.090047	2266.794470
min	0.000000	0.000000	18.250000	0.000000
25%	0.000000	9.000000	35.500000	398.550000
50%	0.000000	29.000000	70.350000	1394.550000
75%	0.000000	55.000000	89.850000	3786.600000
max	1.000000	72.000000	118.750000	8684.800000

Numerical Features - Analysis

Understand the distribution of the numerical features

```
from seaborn.palettes import color_palette
def plot_histogram(df, column_name):

    plt.figure(figsize=(5,3))
    sns.histplot(data=df, x=column_name, kde=True)
    plt.title(f"Distribution of {column_name}")


    # calculate the mean and median values for the columns
    col_mean = df[column_name].mean()
    col_median = df[column_name].median()

    # add vertical lines for mean and median
    plt.axvline(col_mean, color="red", linestyle="--", label="Mean")
    plt.axvline(col_median, color="green", linestyle="-", label="Median")


    plt.legend()

    plt.show()
```


```
plot_histogram(df, "tenure")
```



```
plot_histogram(df, "MonthlyCharges")
```




```
plot_histogram(df, "TotalCharges")
```




Box plot for numerical features

```
def plot_boxplot(df, column_name):  
  
    plt.figure(figsize=(5,3))  
    sns.boxplot(y=df[column_name])  
    plt.title(f"Box Plot of {column_name}")  
    plt.ylabel(column_name)  
    plt.show
```


```
plot_boxplot(df, "tenure")
```



```
plot_boxplot(df, "MonthlyCharges")
```




```
plot_boxplot(df, "TotalCharges")
```



Correlation Heatmap for numerical columns

```
# correlation matrix - heatmap
plt.figure(figsize=(8, 4))
sns.heatmap(df[["tenure", "MonthlyCharges", "TotalCharges"]].corr())
plt.show()
```



Categorical features - Analysis

```
df.columns
```

```
Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents',
       'tenure',
       'PhoneService', 'MultipleLines', 'InternetService',
       'OnlineSecurity',
       'OnlineBackup', 'DeviceProtection', 'TechSupport',
       'StreamingTV',
       'StreamingMovies', 'Contract', 'PaperlessBilling',
       'PaymentMethod',
       'MonthlyCharges', 'TotalCharges', 'Churn'],
       dtype='object')
```

```
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   gender          7043 non-null    object 
 1   SeniorCitizen   7043 non-null    int64  
 2   Partner         7043 non-null    object 
 3   Dependents     7043 non-null    object 
 4   tenure          7043 non-null    int64  
 5   PhoneService    7043 non-null    object 
 6   MultipleLines   7043 non-null    object 
 7   InternetService 7043 non-null    object 
 8   OnlineSecurity  7043 non-null    object 
 9   OnlineBackup    7043 non-null    object 
 10  DeviceProtection 7043 non-null    object 
 11  TechSupport    7043 non-null    object 
 12  StreamingTV    7043 non-null    object 
 13  StreamingMovies 7043 non-null    object 
 14  Contract        7043 non-null    object 
 15  PaperlessBilling 7043 non-null    object 
 16  PaymentMethod   7043 non-null    object 
 17  MonthlyCharges  7043 non-null    float64
 18  TotalCharges    7043 non-null    float64
 19  Churn           7043 non-null    object 
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB
```





Countplot for categorical columns




```
object_cols = df.select_dtypes(include="object").columns.to_list()
```




```
object_cols = ["SeniorCitizen"] + object_cols

for col in object_cols:
    plt.figure(figsize=(5,3))
    sns.countplot(x=df[col])
    plt.title(f"Count Plot of {col}")
    plt.show()
```















Count Plot of StreamingTV





	No	Yes
Churn		

4. Data Preprocessing

```
df.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	...
0	Female	0	Yes	No	1	No	...
1	Male	0	No	No	34	Yes	...
2	Male	0	No	No	2	Yes	...
3	Male	0	No	No	45	No	...
4	Female	0	No	No	2	Yes	...

Next steps:

[Generate code with df](#)[New interactive sheet](#)

Label encoding of target column

```
df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})
```

```
/tmp/ipython-input-2364848822.py:1: FutureWarning: Downcasting behavior
df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})
```

```
df.head(3)
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceFees	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	Churn
0	Female	0	Yes	No	1	No	No	No	No	No	No	No	No	No	No	No	No	No
1	Male	0	No	No	34	No	No	No	No	No	No	No	No	No	No	No	No	Yes
2	Male	0	No	No	2	No	No	No	No	No	No	No	No	No	No	No	No	Yes

Next steps:

[Generate code with df](#)[New interactive sheet](#)

```
print(df["Churn"].value_counts())
```

```
Churn
0    5174
1    1869
Name: count, dtype: int64
```

Label encoding of categorical features

```
# identify columns with object data type
object_columns = df.select_dtypes(include="object").columns
```

```
print(object_columns)
```

```
Index(['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines',
       'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceFees',
       'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
       'PaperlessBilling', 'PaymentMethod'],
      dtype='object')
```

```
# initialize a dictionary to save the encoders
encoders = {}

# apply label encoding and store the encoders
for column in object_columns:
    label_encoder = LabelEncoder()
    df[column] = label_encoder.fit_transform(df[column])
    encoders[column] = label_encoder

#save the encoders to a pickle file
with open("encoders.pkl", "wb") as f:
    pickle.dump(encoders, f)
```

encoders

```
{'gender': LabelEncoder(),
 'Partner': LabelEncoder(),
 'Dependents': LabelEncoder(),
 'PhoneService': LabelEncoder(),
 'MultipleLines': LabelEncoder(),
 'InternetService': LabelEncoder(),
 'OnlineSecurity': LabelEncoder(),
 'OnlineBackup': LabelEncoder(),
 'DeviceProtection': LabelEncoder(),
 'TechSupport': LabelEncoder(),
 'StreamingTV': LabelEncoder(),
 'StreamingMovies': LabelEncoder(),
 'Contract': LabelEncoder(),
 'PaperlessBilling': LabelEncoder(),
 'PaymentMethod': LabelEncoder()}
```

```
df.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	...
0	0		1	0	1	1	0
1	1		0	0	0	34	1
2	1		0	0	0	2	1
3	1		0	0	0	45	0
4	0		0	0	0	2	1

Next steps:

[Generate code with df](#)[New interactive sheet](#)

Training and test data split

```
# splitting the features and target
X = df.drop(columns=["Churn"])
y = df["Churn"]
```

```
# split training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
print(y_train.shape)
```

```
(5634,)
```

```
print(y_train.value_counts())
```

```
Churn
0    4138
1    1496
Name: count, dtype: int64
```

Synthetic Minority Oversampling TEchnique (SMOTE)

```
smote = SMOTE(random_state=42)
```

```
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
```

```
print(y_train_smote.shape)
```

```
(8276,)
```

```
print(y_train_smote.value_counts())
```

```
Churn
0    4138
1    4138
Name: count, dtype: int64
```

5. Model Training

Training with default hyperparameters

```
# dictionary of models
models = {
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "XGBoost": XGBClassifier(random_state=42, use_label_encoder=False),
    "Logistic Regression": LogisticRegression(random_state=42, max_iter=100),
    "SVC": SVC(random_state=42, probability=True)    # probability=True
}
```

```
# dictionary to store the cross validation results
cv_scores = {}

print(" Performing 5-Fold Cross Validation for All Models \n")

# perform 5-fold cross validation for each model
for model_name, model in models.items():
    print(f"Training {model_name} with default parameters...")
    scores = cross_val_score(model, X_train_smote, y_train_smote, cv=5)
    cv_scores[model_name] = scores
    print(f"{model_name} cross-validation accuracy: {np.mean(scores)}")
    print("-" * 70)
```

Performing 5-Fold Cross Validation for All Models

Training Decision Tree with default parameters...
Decision Tree cross-validation accuracy: 0.7778

Training Random Forest with default parameters...
Random Forest cross-validation accuracy: 0.8408

Training XGBoost with default parameters...
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: Parameters: { "use_label_encoder" } are not used.

 bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: Parameters: { "use_label_encoder" } are not used.

 bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: Parameters: { "use_label_encoder" } are not used.

 bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: Parameters: { "use_label_encoder" } are not used.

 bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: Parameters: { "use_label_encoder" } are not used.

 bst.update(dtrain, iteration=i, fobj=obj)
XGBoost cross-validation accuracy: 0.8310

Training Logistic Regression with default parameters...
Logistic Regression cross-validation accuracy: 0.7927

Training SVC with default parameters...
SVC cross-validation accuracy: 0.6396

```
cv_scores
```

```
{'Decision Tree': array([0.68297101, 0.71299094, 0.82175227,
 0.83564955, 0.83564955]),
 'Random Forest': array([0.72524155, 0.77824773, 0.90513595,
 0.89425982, 0.90090634]),
 'XGBoost': array([0.70048309, 0.75649547, 0.90271903, 0.89486405,
 0.90030211]),
 'Logistic Regression': array([0.73188406, 0.74803625, 0.8265861 ,
 0.81993958, 0.83685801]),
 'SVC': array([0.65519324, 0.65740181, 0.61510574, 0.61993958,
 0.65015106])}
```

6. Model Evaluation

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
evaluation_results = {}

print("\n Evaluating Models on Test Data \n")

# Loop through all 5 models and evaluate each
for model_name, model in models.items():
    print(f"Evaluating {model_name}...")
    model.fit(X_train_smote, y_train_smote)      # train each model
    y_pred = model.predict(X_test)                 # test predictions

    acc = accuracy_score(y_test, y_pred)
    cm = confusion_matrix(y_test, y_pred)
    cr = classification_report(y_test, y_pred, output_dict=True)

    evaluation_results[model_name] = {
        "Accuracy": acc,
        "Precision": cr['1']['precision'],
        "Recall": cr['1']['recall'],
        "F1-Score": cr['1']['f1-score']
    }

    print(f"Accuracy: {acc:.4f}")
    print("Confusion Matrix:\n", cm)
    print("Classification Report:\n", classification_report(y_test, y_pred))
    print("-" * 70)

# Create a summary DataFrame to compare model performance
eval_df = pd.DataFrame(evaluation_results).T
print("\n Model Performance Comparison:")
display(eval_df.sort_values(by="Accuracy", ascending=False))
```

Evaluating Models on Test Data

Evaluating Decision Tree...

Accuracy: 0.7317

Confusion Matrix:

```
[[824 212]
 [166 207]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.80	0.81	1036
1	0.49	0.55	0.52	373
accuracy			0.73	1409
macro avg	0.66	0.68	0.67	1409
weighted avg	0.74	0.73	0.74	1409

Evaluating Random Forest...

Accuracy: 0.7786

Confusion Matrix:

```
[[878 158]
 [154 219]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.85	0.85	1036
1	0.58	0.59	0.58	373
accuracy			0.78	1409
macro avg	0.72	0.72	0.72	1409
weighted avg	0.78	0.78	0.78	1409

Evaluating XGBoost...

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: Use Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

Accuracy: 0.7807

Confusion Matrix:

```
[[885 151]
 [158 215]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.85	0.85	1036
1	0.59	0.58	0.58	373
accuracy			0.78	1409
macro avg	0.72	0.72	0.72	1409
weighted avg	0.78	0.78	0.78	1409

weighted avg	0.78	0.78	0.78	1409
--------------	------	------	------	------

Evaluating Logistic Regression...

Accuracy: 0.7644

Confusion Matrix:

```
[[784 252]
 [ 80 293]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.76	0.83	1036
1	0.54	0.79	0.64	373
accuracy			0.76	1409
macro avg	0.72	0.77	0.73	1409
weighted avg	0.81	0.76	0.78	1409

Evaluating SVC...

Accuracy: 0.6884

Confusion Matrix:

```
[[733 303]
 [136 237]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.71	0.77	1036
1	0.44	0.64	0.52	373
accuracy			0.69	1409
macro avg	0.64	0.67	0.64	1409
weighted avg	0.74	0.69	0.70	1409

Model Performance Comparison:

	Accuracy	Precision	Recall	F1-Score	
XGBoost	0.780696	0.587432	0.576408	0.581867	
Random Forest	0.778566	0.580902	0.587131	0.584000	
Logistic Regression	0.764372	0.537615	0.785523	0.638344	
Decision Tree	0.731725	0.494033	0.554960	0.522727	
SVC	0.688432	0.438889	0.635389	0.519168	

7.Selecting and Saving the Best Model

Save the best model

```
best_model_name = eval_df["Accuracy"].idxmax()
best_model = models[best_model_name]

print(f"\n Best Model: {best_model_name} with Accuracy = {eval_df.loc[best_model_name, 'Accuracy']}")

best_model.fit(X_train_smote, y_train_smote)

model_data = {"model": best_model, "feature_names": X.columns.to_list()}

import pickle
with open("customer_churn_best_model.pkl", "wb") as f:
    pickle.dump(model_data, f)

print(" Best model saved as 'customer_churn_best_model.pkl'")
```

```
Best Model: XGBoost with Accuracy = 0.7807
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
Best model saved as 'customer_churn_best_model.pkl'
```

Load the saved best model

```
with open("customer_churn_best_model.pkl", "rb") as f:  
    model_data = pickle.load(f)  
  
loaded_model = model_data["model"]  
feature_names = model_data["feature_names"]  
  
print(" Loaded model:", loaded_model)
```

```
Loaded model: XGBClassifier(base_score=None, booster=None, callback=  
    colsample_bylevel=None, colsample_bynode=None,  
    colsample_bytree=None, device=None, early_stopping_rounds=None,  
    enable_categorical=False, eval_metric='logloss',  
    feature_types=None, feature_weights=None, gamma=None,  
    grow_policy=None, importance_type=None,  
    interaction_constraints=None, learning_rate=None, max_depth=None,  
    max_cat_threshold=None, max_cat_to_onehot=None,  
    max_delta_step=None, max_depth=None, max_leaves=None,  
    min_child_weight=None, missing=nan, monotone_constraints=None,  
    multi_strategy=None, n_estimators=None, n_jobs=None,  
    num_parallel_tree=None, ...)
```

Predict Churn for a Sample Customer (Quick Test)... Static Datas

```
import pandas as pd

input_data = {
    "gender": "Female",
    "SeniorCitizen": 0,
    "Partner": "Yes",
    "Dependents": "No",
    "tenure": 1,
    "PhoneService": "No",
    "MultipleLines": "No phone service",
    "InternetService": "DSL",
    "OnlineSecurity": "No",
    "OnlineBackup": "Yes",
    "DeviceProtection": "No",
    "TechSupport": "No",
    "StreamingTV": "No",
    "StreamingMovies": "No",
    "Contract": "Month-to-month",
    "PaperlessBilling": "Yes",
    "PaymentMethod": "Electronic check",
    "MonthlyCharges": 29.85,
    "TotalCharges": 29.85
}

# Load encoders
with open("encoders.pkl", "rb") as f:
    encoders = pickle.load(f)

input_df = pd.DataFrame([input_data])

# Encode categorical features
for col, encoder in encoders.items():
    input_df[col] = encoder.transform(input_df[col])

# Predict
pred = loaded_model.predict(input_df)
prob = loaded_model.predict_proba(input_df)[0][1]

print(" Prediction Result:")
print(f" Prediction: {'Churn' if pred[0]==1 else 'No Churn'}")
print(f" Churn Probability: {prob:.2f}")
```

